

(ADAM) CHEOL WOO KIM

(+1) 917-985-1869 $\diamond$ cwkim@seas.harvard.edu $\diamond$ acwkim.github.io $\diamond$ Updated Sep 2026

ACADEMIC APPOINTMENTS

Harvard University

Postdoctoral Fellow in Computer Science

- Advisor: Milind Tambe

Boston, MA

Sep 2024 – Present

EDUCATION

Massachusetts Institute of Technology

PhD in Operations Research; GPA: 5.0/5.0

- Advisor: Dimitris Bertsimas

- Thesis: *Predictive and Prescriptive Trees for Optimization and Control Problems*

Cambridge, MA

Aug 2024

Seoul National University School of Law

J.D. Candidate (withdrew to pursue PhD at MIT)

Seoul, Republic of Korea

Feb 2019 – Jul 2019

Seoul National University

B.A. in Economics (Summa Cum Laude); Major GPA: 4.24/4.30

- Concentration in Mathematical Economics and Microeconomic Theory

Seoul, Republic of Korea

Feb 2019

RESEARCH INTERESTS

AI Alignment and Personalization, AI Safety, Human-AI Interaction,
Learning to Optimize, AI-Assisted Decision Modeling

PAPERS *: equal contribution; α - β : alphabetical order

Preprints

1. Many Preferences, Few Policies: Compact Portfolios for Multi-objective LLM Alignment
Cheol Woo Kim*, Jai Moondra*, Roozbeh Nahavandi, Andrew Perrault, Milind Tambe, Swati Gupta
Workshop version: *Conference on Language Modeling (COLM), Workshop on Scientific Understanding of Foundation Models*, 2026.

Journal Papers (Including Under Revision)

1. Optimal Control of Multiclass Fluid Queueing Networks: A Machine Learning Approach
(α - β) Dimitris Bertsimas, **Cheol Woo Kim**
Operations Research (minor revision), 2026.
2. Optimal Control of Fluid Restless Multi-armed Bandits: A Machine Learning Approach
(α - β) Dimitris Bertsimas, **Cheol Woo Kim**, José Niño-Mora
Machine Learning, 2026.
3. A Machine Learning Approach to Two-stage Adaptive Robust Optimization
(α - β) Dimitris Bertsimas, **Cheol Woo Kim**
European Journal of Operational Research, 2024.
4. A Prescriptive Machine Learning Approach to Mixed-Integer Convex Optimization
(α - β) Dimitris Bertsimas, **Cheol Woo Kim**
INFORMS Journal on Computing, 2023.

Conference Papers

1. Bilevel Optimization of Synthetic Trajectories for Multi-Turn LLM Fine-Tuning
Shresth Verma, Mauricio Tec, **Cheol Woo Kim**, Kai Wang, Milind Tambe
Advances in Neural Information Processing Systems (NeurIPS), 2026.
2. Incentive-Aware AI Safety via Strategic Resource Allocation: A Stackelberg Security Games Perspective
Cheol Woo Kim*, Davin Choo*, Tzeh Yuan Neoh, Milind Tambe
International Conference on Autonomous Agents and Multiagent Systems (AAMAS), Blue Sky Ideas Track, 2026.
3. Preference Robustness for DPO with Applications to Public Health
Cheol Woo Kim*, Shresth Verma*, Mauricio Tec*, Milind Tambe
AAAI Conference on Artificial Intelligence (AAAI), AI for Social Impact Track, 2026.
4. Navigating the Social Welfare Frontier: Portfolios for Multi-objective Reinforcement Learning
Cheol Woo Kim*, Jai Moondra*, Shresth Verma, Madeleine Pollack, Ling kai Kong, Milind Tambe, Swati Gupta
International Conference on Machine Learning (ICML), 2025.
5. Robust Optimization with Diffusion Models for Green Security
Ling kai Kong, Haichuan Wang, Yuqi Pan, **Cheol Woo Kim**, Mingxiao Song, Alayna Nguyen, Tonghan Wang, Haifeng Xu, Milind Tambe
Uncertainty in Artificial Intelligence (UAI), 2025.

Workshop Papers

1. Lightweight Robust Direct Preference Optimization
Cheol Woo Kim*, Shresth Verma*, Mauricio Tec*, Milind Tambe
NeurIPS Workshop: Reliable ML from Unreliable Data, 2025.
2. LLM-based Agent Simulation for Maternal Health Interventions: Uncertainty Estimation and Decision-focused Evaluation
Sarah Martinson, Ling kai Kong, **Cheol Woo Kim**, Aparna Taneja, Milind Tambe
AAMAS Workshop on Autonomous Agents for Social Good, 2025.

WORK EXPERIENCE

Microsoft Research	Redmond, WA
<i>Research Intern, Machine Learning and Optimization Group</i>	<i>May 2023 – Aug 2023</i>
• Mentors: Ishai Menache, Marco Molinaro, Konstantina Mellou	
• Research: Combinatorial optimization under uncertainty for GPU resource allocation.	
Permanent Mission of the Republic of Korea to the United Nations	New York, NY
<i>Research Intern, UN General Assembly Group</i>	<i>Jan 2014 – Apr 2014</i>

TEACHING EXPERIENCE

Developed course materials, led recitations, and held office hours for the following courses:

- TA, The Analytics Edge (MIT 15.071)** *Spring 2024*
• MBA course on statistics, machine learning, optimization, and data science.
- TA, Optimization Methods (MIT 6.7200J/15.093J)** *Fall 2022*
• Master's/PhD course on linear, nonlinear, combinatorial, robust, and semi-definite optimization.
- TA, Introduction to Mathematical Programming (MIT 6.251J/15.081J)** *Fall 2021*
• PhD course on linear, nonlinear, discrete, combinatorial, robust, and semi-definite optimization.

TA, Robust Modeling, Optimization, and Computation (MIT 1.142J/15.094J) *Spring 2021*

- Master's/PhD course on theory and applications of robust optimization.

AWARDS AND HONORS

Kwanjeong Educational Fellowship

Sep 2019 – Feb 2025

- Merit-based PhD scholarship.

Eminence Scholarship

Aug 2017

Department of Economics, Seoul National University

Seoul, Republic of Korea

- Awarded to top 1% in the economics department.

INVITED TALKS

Workshop on Learning Algorithms for Decision-Making in Games with Mixed Agents, Cornell University

Apr 2026

- *Towards Scalable Language Model Personalization*

KAIST Graduate School of Data Science Special Seminar

Aug 2024

- *Machine Learning for Optimization and Control Problems*

INFORMS Annual Meeting

Oct 2023

- *Optimal Control of Fluid Restless Multi-armed Bandits: A Machine Learning Approach*

PROFESSIONAL SERVICE

Journal Reviewer: Journal of Machine Learning Research, INFORMS Journal on Computing, INFORMS Journal on Optimization, European Journal of Operational Research, Transportation Science

Conference Reviewer: AAAI 2026

GRANT PROPOSAL CONTRIBUTIONS

Contributed substantially to writing the following funded proposal:

- **Interactive Strategic Policy Design using LLM Agents**

Army Research Office; PI: Milind Tambe

Aug 2026 – Jul 2029

STUDENTS ADVISED

Robert Debbas, Harvard University

2024 – Present

Shresth Verma, Harvard University

2024 – Present

Sarah Martinson, Harvard University

2024 – 2025

Next position: Ph.D. student in Applied Mathematics, Harvard University.

SKILLS

Programming: Python (PyTorch, scikit-learn), Julia, R

Optimization Software: JuMP, Gurobi, MOSEK

Languages: Korean (native), English (fluent)